

Confidence intervals

Let's look at the Gallup poll on the approval rating of the US President. Suppose 60% of the 140 million likely voters approve of the way the president is handling his job.

Gallup polls 1,000 of them. The resulting approval percentage in the sample will be off the population percentage of 60% due to chance error. How much?

The SE tells us the likely size of the chance error.

Confidence intervals give a more precise statement.

Confidence intervals

According to the central limit theorem, the sample percentage follows the normal curve with expected value $\mu = 60\%$ and SE equal to $\frac{\sigma}{\sqrt{1000}} = \frac{0.49}{31.6} = 1.6\%$.

(Because we sample from a population of 140 million labels, of which 60% are 1s and 40% are 0s.)

By the empirical rule there is a 95% chance that the sample percentage is no more than 2 SEs away from the population percentage μ .

Confidence intervals

By the empirical rule there is a 95% chance that the sample percentage is no more than 2 SEs away from the population percentage μ .

But that is the same as saying that the population percentage is no more than 2 SEs away from the sample percentage.

So once we compute the sample percentage from our sample, say 58%, then we can give a range of plausible values for the unknown population percentage by going 2 SEs in each direction:

[54.8%, 61.2%] is a 95% **confidence interval** for the population percentage.

Interpretation of a confidence interval

[54.8%, 61.2%] is a 95% confidence interval for the population percentage.

Why 'confidence' instead of 'probability'?

The population percentage μ is a *fixed* number, which either falls into [54.8%, 61.2%] or not. There are no chances involved.


Rather, the chances are in the sampling procedure:

a different sample of 1,000 voters will give a slightly different interval.

If one does many polls (the Gallup poll is done frequently), then 95% of these intervals trap the population percentage, and 5% will miss it.

95% is called the **confidence level**.

Interpretation of a confidence interval



"I am 95% confident that the President's approval rating is between 54.8% and 61.2%" means that 95% of the time I am correct when making such a statement based on a poll.

Keep in mind that the interval varies from sample to sample, while the population percentage is a fixed number.

Confidence intervals via the Central Limit Theorem

A confidence interval gives a range of plausible values for a *population parameter* μ . Usually the confidence interval is centered at an estimate for μ which is an average.

Examples:

1. μ = approval percentage among all 140 million likely voters.
estimate = approval percentage among voters in the sample.
2. μ = speed of light.
estimate = average of 30 measurements.

Since the central limit theorem applies for averages, the confidence interval has a simple form:

$$\text{estimate} \pm z \text{ SE}$$

where z is the z-score corresponding to the desired confidence level:

95% confidence level $\rightarrow z = 1.96$

90% confidence level $\rightarrow z = 1.65$

99% confidence level $\rightarrow z = 2.58$

Estimating the SE with the bootstrap principle

SE is the standard error of the estimate. If the estimate is an average (e.g. a percentage), then we know that

$$\text{SE} = \frac{\sigma}{\sqrt{n}}$$

where σ is the standard deviation of the population.

Now we have a problem: we don't know σ because we don't know all the data in the population (that's the reason why we sample in the first place!)

The **bootstrap principle** states that we can estimate σ by its sample version s and still get an approximately correct confidence interval.

Estimating the SE with the bootstrap principle

The **bootstrap principle** states that we can estimate σ by its sample version s and still get an approximately correct confidence interval.

Examples:

1. We poll 1,000 likely voters and find that 58% approve of the way the president handles his job.

$SE = \frac{\sigma}{\sqrt{n}} \times 100\%$, where $\sigma = \sqrt{p(1-p)}$, p = proportion of all voters who approve.

The bootstrap principle replaces σ by s = standard deviation of the 0/1 labels in the sample = $\sqrt{0.58(1-0.58)} = 0.49$.

So a 95% confidence interval for p is

$$58\% \pm 2 \frac{0.49}{\sqrt{1000}}, \quad \text{which is } [54.9\%, 61.1\%]$$

Estimating the SE with the bootstrap principle

2. estimate = average of 30 measurements of the speed of light.

What is the population or probability histogram from which we sample?

The reason we get 30 different measurementss is because each is off by a chance error:

$$\text{measurement} = \text{speed of light} + \text{measurement error}$$

The measurement error follows a probability histogram that is unknown to us. We estimate the standard deviation σ of this probability histogram by the standard deviation s of the sample of 30 measurements.

Later we will see how the bootstrap principle can be used to approximate the SE of estimates other than averages.

More about confidence intervals

- ▶ The width of the confidence interval is determined by z SE, which is called the **margin of error**.

A larger sample size n will result in a smaller margin of error since $SE = \frac{\sigma}{\sqrt{n}}$, but we need *four times the sample size to cut the width in half*.

We can also make the width smaller by making z smaller, e.g. use a 80% confidence level instead of 95%. Then the price for more precision (i.e. shorter interval) is less confidence that it covers the parameter μ .

- ▶ There is an easy to remember formula for a 95% confidence interval for a percentage:

$$\text{estimated percentage} \pm \frac{1}{\sqrt{n}}$$

That's because $\sigma = \sqrt{p(1-p)} \leq \frac{1}{2}$ no matter what p is.

- ▶ It is journalistic convention to use a 95% confidence level unless stated otherwise.